



BLG-413E – System Programming

150130066 Duhan Cem Karagöz

150150701 Yunus G ng r

Homework 1

13.10.2017

Objective:

Objective of the homework is to learn how to implement a system call to a linux based system and learn how to develop and change already existing system calls and functions.

System:

On this homework Ubuntu 14.04 with GCC 4.7 is used.

Program:

At this homework we based our code and thoughts on classes and practice session slides. For this homework we are assigned to develop a way to kill a processes children if a flag is set to one and priority of that process is higher than 30. So we added a flag value to task descriptor and check priority whenever process exits.

```
1464  #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
1465      unsigned int    sequential_io;
1466      unsigned int    sequential_io_avg;
1467  #endif
1468      int myFlag;
1469  };
```

Figure 1. sched.h

To be able to create a flag and change it first of all we added a flag to end of the task descriptor. We needed to add our flag to end of the descriptor. We could not add the flag in between since it would create misalignment issues with the jump operation in machine code.

```
212      [PIDTYPE_PID] = INIT_PID_LINK(PIDTYPE_PID),      \
213      [PIDTYPE_PGID] = INIT_PID_LINK(PIDTYPE_PGID),    \
214      [PIDTYPE_SID] = INIT_PID_LINK(PIDTYPE_SID),      \
215  },                                                    \
216  .thread_group = LIST_HEAD_INIT(tsk.thread_group),    \
217  .thread_node = LIST_HEAD_INIT(init_signals.thread_head), \
218  .myFlag=0, \
219  INIT_IDS \
220  INIT_PERF_EVENTS(tsk) \
221  INIT_TRACE_IRQFLAGS \
222  INIT_LOCKDEP \
223  INIT_FTRACE_GRAPH \
224  INIT_TRACE_RECURSION \
225  INIT_TASK_RCU_PREEMPT(tsk) \
226  INIT_CPUSET_SEQ(tsk) \
227  INIT_VTIME(tsk) \
228  }
```

Figure 2. init_task.h

When first process created it will be allocated statically and its task descriptor is taken from `init_task.h` therefore we need to initialize its flag as zero at the file.

```
7  asmlinkage long set_myFlag(pid_t pid, int flag){
8
9      printk("set my flag starts running with: %d, %d \n",pid,flag);
10
11     struct task_struct *task;
12     task=find_task_by_vpid(pid);
13
14     if (task>0)
15     {
16         printk("found task: %d \n", task->pid);
17         if(current_cred()->0){
18             if(current_cred()->uid <= 0 || current_cred()->gid<=0){ //https://www.kernel.org/doc/Documentation/security/credentials.txt
19                 //Process has root priv. check flag
20                 if(flag==1 || flag==0){
21                     task->myFlag=flag;
22                     printk("my_flag value of process %d has been set to %d \n",pid,task->myFlag);
23                     //https://www.ibm.com/developerworks/library/l-kernel-logging-apis/index.html
24                     return 0;
25                 }
26             }
27             else{
28                 printk("my_flag value of process %d can not set to %d \n",pid,flag);
29                 return -EINVAL;
30             }
31         }
32         else{
33             //Not Root throw error
34             printk("my_flag value of process %d can not set to %d \n",pid,flag);
35             //https://www.ibm.com/developerworks/library/l-kernel-logging-apis/index.html
36             return -EPERM; // http://man7.org/linux/man-pages/man2/syscalls.2.html
37         }
38     }
39     else{
40         printk("user error \n");
41         return -EACCES;
42     }
43 }
44 else{
45     printk("no process found \n");
46     return -ESRCH;
47 }
48 }
49 }
```

Figure 3. `set_myFlag.c`

When a process needs to change its flag value it needs to call this function with a pid number but calling process must have root privileges. At the first if it checks if pid is valid then checks if calling process has root privileges then it will assign the new flag value to process. If a check is failed then corresponding error will be returned to the calling process.

```
847  asmlinkage long sys_kcmp(pid_t pid1, pid_t pid2, int type,
848      unsigned long idx1, unsigned long idx2);
849  asmlinkage long sys_finit_module(int fd, const char __user *uargs, int flags);
850  asmlinkage long sys_seccomp(unsigned int op, unsigned int flags,
851      const char __user *uargs);
852  asmlinkage long set_myFlag(pid_t pid, int flag);
853  #endif
854
```

Figure 4. `sys_calls.c`

We added a prototype of our call to `sys_calls` to be recognized and compiled.

```

359 350 i386   finit_module      sys_finit_module
360 351 i386   sched_setattr     sys_ni_syscall
361 352 i386   sched_getattr     sys_ni_syscall
362 353 i386   renameat2         sys_ni_syscall
363 354 i386   seccomp           sys_seccomp
364 355 i386   set_myFlag       set_myFlag
365

```

Figure 5. syscall_32.tbl

To use our new system call we added the system call to syscall_32.tbl.

```

344     atomic_set(&tsk->usage, 2);
345     #ifdef CONFIG_BLK_DEV_IO_TRACE
346     tsk->btrace_seq = 0;
347     #endif
348     tsk->splice_pipe = NULL;
349     tsk->task_frag.page = NULL;
350
351     account_kernel_stack(ti, 1);
352
353     tsk->myFlag = 0 ;
354
355     return tsk;
356
357 free_ti:
358     free_thread_info(ti);
359 free_tsk:
360     free_task_struct(tsk);
361     return NULL;
362 }

```

Figure 6. fork.c

When a new process is forked child process copies its task from its parent. So at every fork we set flag that copied to zero to prevent any mishaps from happening.

```

708     struct task_struct *tsk = current;
709     int group_dead;
710
711     struct task_struct *task;
712     struct list_head *list;
713     int pids;
714     int priority;
715     int tsk_flag=tsk->myFlag;
716     priority=sys_getpriority(PRIO_PROCESS, tsk->pid);
717
718     if(tsk_flag == 1 && priority>10)
719     {
720         printk("Parents with pid pid:%d proi is %d.\n",tsk->pid, priority);
721         printk("Parent with pid:%d has been terminated Killing offsprings.\n",tsk->pid);
722         //get all childs pid and call them to exit
723         list_for_each(list, &current->children)
724         {
725             task = list_entry(list, struct task_struct, sibling);
726             pids = task->pid;
727             printk("Child with pid:%d has been terminated \n",task->pid);
728             sys_kill(task->pid, SIGKILL);
729             // task points to a children
730         }
731         //Linux Kernel Development - 3rd Edition Page 30
732     }

```

Figure 7. exit.c

To be able to kill child processes with parent we altered the exit.c Every time a process exits it calls exit() function that in turn calls do_exit() so we added a code line to first part of do_exit() everytime a process exits it will go through this lines of code. If a process arrives with flag 1 and priority above 10 it will go into the if statement. Since task descriptors are connected via double link list it is easy find child processes. With the list_for_each macro it will find each child of the current process and kill it.

```

2  #define NR_mycall 355
3  #include <unistd.h>
4  #include <sys/time.h>
5  #include <sys/resource.h>
6  #include <stdlib.h>
7
8  int main (void){
9      int y;
10     int pid;
11     int which=PRIO_PROCESS;
12     id_t pids=getpid();
13     int secim=0;
14     printf("1-> Flag With 0 And Prio 0\n");
15     printf("2-> Flag With 1 And Prio 0\n");
16     printf("3-> Flag With 0 And Prio 33\n");
17     printf("4-> Flag With 1 And Prio 33\n");
18     printf("5-> Flag With 2 And Prio 33\n");
19

```

Figure 8. test_set_myValue.c Part 1.

To test our changes in the kernel we wrote a piece of code to check if our implementation is worked correctly. First part is to select one of five different scenario to check our code. First you can check if code is not working when it should not with not setting anything. Second scenario is to check if our priority checking procedures are working. Third is for flag checking. Fourth one is to check if when everything is set is our program working accordingly. Meaning is our program killing childs when it suppose to kill.

```
22     switch(secim)
23     {
24     case 1:
25         //do not set both up leave them down
26         break;
27
28     case 2:
29         y=syscall(NR_mycall, pids, 1);
30         //set flag one but leave prio low
31         break;
32
33     case 3:
34         setpriority(which, pids, -15);
35         //set prio high but flag 0
36         break;
37
38     case 4:
39         y=syscall(NR_mycall, pids, 1);
40         setpriority(which, pids, -15);
41         //set both up
42         break;
43
44     case 5:
45         y=syscall(NR_mycall, pids, 2);
46         setpriority(which, pids, -15);
47         //set both up
48         break;
49
50     default:
51         return 1;
52     }
```

Figure 9. test_set_myValue.c Part 2.

This part shows how we set flag and priorities.

```

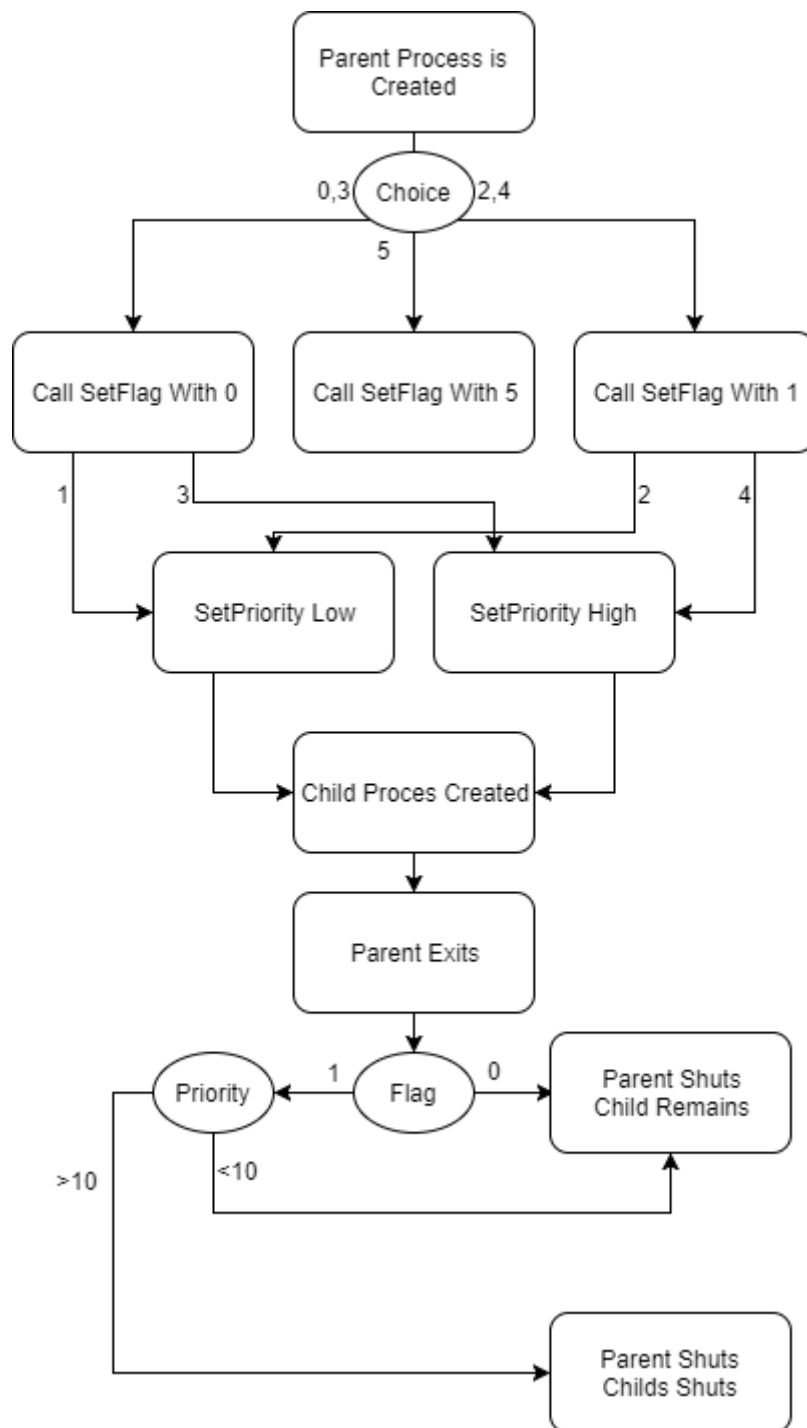
54     pid=fork();
55
56     if(pid==0)
57     {
58         printf("I am the child\n");
59         printf("my pid=%d\n", getpid());
60         while(1)
61         {
62             //empty must wait father to exit
63             //or CTRL+C to stop
64         }
65
66         exit(0);
67     }
68
69     else
70     {
71         int waits;
72         printf("I am the father\n");
73         printf("my pid=%d\n", getpid());
74         printf("My Prio %d\n", getpriority(which, getpid()));
75         //sleep(1);
76         scanf("%d", &waits);
77         exit(0);
78         //return 1;
79     }
80

```

Figure 10. test_set_myValue.c Part 3.

This part is where we forked the process and put child process into infinite loop to see if it exits or not with the parent process.

Flow Diagram:



Results:

We learned how to implement a system call and change already existing ones. Also we learned how to add a value to task descriptor. And we learned how to develop and improve existing kernels.